

Linear Text classification

<https://powcoder.com>

Assignment Project Exam Help

Problem definition: Given a text document, assign it a discrete label $y \in \mathcal{Y}$ where \mathcal{Y} is the set of possible labels. Many possible applications:

- ▶ Spam filter: $\mathcal{Y} = \{\text{Spam, non-spam}\}$
- ▶ Sentiment: $\mathcal{Y} = \{\text{Positive, negative, neutral}\}$
- ▶ Genre classification: $\mathcal{Y} = \{\text{sports, fiction, news, } \dots\}$

Bag-of-words representation of a document

<https://powcoder.com>

A typical representation of a document is a bag of words, which is mathematically a vector of word counts:

Assignment Project Exam Help

Assignment Project Exam Help

Example $x = [1, 5, 3, 3, 10, 0, 1]$

where x_j is the count of the word j . The length of the vector is the size of the vocabulary $|V|$. (So that the vector for all documents in a set are of the same length for apple-to-apple comparison.)

<https://powcoder.com>

Add WeChat powcoder

- ▶ “Vocabulary” here can’t be understood as words in a language. For instance, it could include all bigrams in a collection of documents.
- ▶ Alternatives to word counts include simple presence 1 or absence 0 of a word, tf/idf of a word, etc.
- ▶ By using word count we dropped all word order information

Feature function

- ▶ Not all words are equally important for purposes of predicting a particular label. To predict the label of a document, we assign a score to each word in the vocabulary to indicate the “compatibility” with the label, e.g., “basketball” has a high compatibility with *sports*, “Gryffindor” has a high compatibility with *fiction*.
- ▶ These compatibility scores are called *weights* and they are arranged in a vector θ .
- ▶ Given a bag-of-words \mathbf{x} and a weight vector θ , we predict the label y by computing the total compatibility score between \mathbf{x} and y .
- ▶ In a linear function, this compatibility score is the inner product of between the weights θ and a *feature function*:

$$\hat{y} = \operatorname{argmax}_{y \in \mathcal{Y}} \Psi(\mathbf{x}, y)$$

$$\Psi(\mathbf{x}, y) = \theta \cdot \mathbf{f}(\mathbf{x}, y) = \sum_i \theta_i f_i(\mathbf{x}, y)$$

More on feature function

<https://powcoder.com>

- ▶ The feature function has two arguments, the word counts \mathbf{x} and the label y .
- ▶ It will return a feature vector where each element of the vector might be:

$$f_j(\mathbf{x}, y) = \begin{cases} x_{whale}, & \text{if } y = \text{Fiction} \\ 0, & \text{Otherwise} \end{cases}$$

- ▶ In this case the size of the feature vector is the size of the vocabulary, but it doesn't have to be.
- ▶ The output of the feature function also doesn't have to be the word count.

Shape of the feature vector

<https://powcoder.com>

$$f(x, y = 1) = [x; \underbrace{0; 0; \dots; 0}_{(K-1) \times V}]$$

$$f(x, y = 2) = [\underbrace{0; 0; \dots; 0}_V; \underbrace{0; 0; \dots; 0}_{(K-2) \times V}; x]$$

$$f(x, y = K) = [\underbrace{0; 0; \dots; 0}_{(K-1) \times V}; x]$$

Add WeChat powcoder

where K is size of the label set, $\underbrace{0; 0; \dots; 0}_{(K-1) \times V}$ is a vector of

$(K - 1) \times V$ zeros and the semicolon indicates vertical concatenation.

Note: Think of a feature vector this way is good for mathematical presentation. It doesn't have to be implemented this way.

Bias

It is common to add a *offset feature* or *bias* at the end of the vector of word counts which is always 1, and then we need to add a zero to each of the zero vectors, to make the vector lengths match. The entire vector $f(\mathbf{x}, y)$ will be $(V + 1)K$

$$f(\mathbf{x}, y = 1) = [\mathbf{x}; \underbrace{0; 0; \dots; 0}_{(K-1) \times (V+1)}]$$

$$f(\mathbf{x}, y = 2) = [\underbrace{0; 0; \dots; 0}_V; \mathbf{x}; \underbrace{0; 0; \dots; 0}_{(K-2) \times (V+1)}]$$

$$f(\mathbf{x}, y = K) = [\underbrace{0; 0; \dots; 0}_{(K-1) \times (V+1)}; \mathbf{x}]$$

Bias: What's its effect on classification if it is the only feature?

Example “vocabulary” V

<https://powcoder.com>

Assignment Project Exam Help

Vocabulary for a collection of documents

A	NEG	not funny at all
B	NEG	painful, not funny
C	NEU	ok overall
D	POS	funny story
E	POS	good story, good jokes

$V = \{ \text{not, funny, painful, ok, overall, story, good, jokes} \}$

From vocabulary to feature function

<https://powcoder.com>

Vocabulary for a collection of documents

Assignment Project Exam Help

A	NEG	not funny at all
B	NEG	not funny at all
C	NEU	ok overall
D	POS	funny story
E	POS	good story, good jokes

Assignment Project Exam Help

<https://powcoder.com>

[Add WeChat powcoder](#)
 $V = \{ \text{not, funny, painful, ok, overall, story, good, jokes} \}$

$$f_1(\mathbf{x}, y) = \begin{cases} x_{not}, & \text{if } y = \text{NEG} \\ 0, & \text{Otherwise} \end{cases}$$

From vocabulary to feature function

<https://powcoder.com>

Vocabulary for a collection of documents

A	NEG	not funny at all
B	NEG	not funny, painful
C	NEU	ok overall
D	POS	funny story
E	POS	good story, good jokes

$V = \{ \text{not, funny, painful, ok, overall, story, good, jokes} \}$

$$f_2(\mathbf{x}, y) = \begin{cases} x_{\text{funny}}, & \text{if } y = \text{NEG} \\ 0, & \text{Otherwise} \end{cases}$$

From vocabulary to feature function

<https://powcoder.com>

Vocabulary for a collection of documents

Assignment Project Exam Help

Assignment Project Exam Help

A	NEG	not funny at all
B	NEG	painful, not funny
C	NEG	ok overall
D	POS	funny story
E	POS	good story, good jokes

<https://powcoder.com>

$V = \{ \text{not, funny, painful, ok, overall, story, good, jokes} \}$

Feature vector for A:

Add WeChat powcoder

$$f(\mathbf{x} = \text{featurize}(A), y = \text{NEG}) = [1; 1; 0; 0; 0; 0; 0; 0; 1; \underbrace{0; 0; \dots; 0}_{(3-1) \times (8+1)}]$$

From vocabulary to feature function

<https://powcoder.com>
Vocabulary for a collection of documents

A	NEG	not funny at all
B	NEG	painful, not funny
C	NEU	ok overall
D	POS	funny story
E	POS	good story, good jokes

$V = \{ \text{not, funny, painful, ok, overall, story, good, jokes} \}$

Feature vector for C

$$\mathbf{f}(\mathbf{x} = \text{featurize}(C), y = NEU) =$$
$$\underbrace{[0; 0; \cdots; 0; 0; 0; 0; 1; 1; 0; 0; 0; 1; 0; 0; \cdots; 0]}_{(8+1)} \underbrace{[0; 0; \cdots; 0]}_{(8+1)}$$

From vocabulary to feature function

<https://powcoder.com>
Vocabulary for a collection of documents

A	NEG	not funny at all
B	NEG	painful, not funny
C	NEU	ok overall
D	POS	funny story
E	POS	good story, good jokes

$V = \{ \text{not, funny, painful, ok, overall, story, good, jokes} \}$

Feature vector for E

$$\begin{aligned} f(\mathbf{x} = \text{featurize}(E), y = POS) = \\ [0; 0; \cdots; 0; 0; 0; 0; 0; 0; 0; 1; 1; 1] \\ \underbrace{\hspace{1.5cm}}_{(3-1) \times (8+1)} \end{aligned}$$

The importance of feature functions

<https://powcoder.com>

Assignment Project Exam Help

- ▶ The performance of a model to a large extent depends on the use of proper feature functions
- ▶ A lot of research went to find the most effective features when developing machine learning systems
- ▶ Models that can't handle a large number of features usually don't perform as well
- ▶ The most effective features differ from task to task, and hence relies on a good understanding of the problem at hand and domain knowledge

Assigning weights to features

<https://powcoder.com>

Assignment Project Exam Help

Now we know about features. What about the weights (θ)?

Assignment Project Exam Help

$$\Psi(\mathbf{x}, y) = \mathbf{f}(\mathbf{x}, y)\theta$$

<https://powcoder.com>

- ▶ There are many different ways to estimate the weights θ . That's why we have different machine learning models.
- ▶ We find the optimal value of θ with a set of training samples of size N : $\{\mathbf{x}^{1:N}, y^{1:N}\}$

Probability preliminaries

<https://powcoder.com>

- ▶ Joint probability: $P(X = a, Y = b)$ or $P(a, b)$ where X and Y are random variables and a and b are values assigned to the random variables.
- ▶ Conditional probability: $P(X = a|Y = b) = \frac{P(X=a, Y=b)}{P(Y=b)}$
- ▶ Bayes' theorem: $P(X = a|Y = b) = \frac{P(Y=b|X=a)P(X=a)}{P(Y=b)}$
- ▶ Marginalization (sum rule): $P(Y) = \sum_x P(X, Y)$
- ▶ Independence: $P(Y|X) = P(Y)$, $P(X|Y) = P(X)$, $P(X, Y) = P(X)P(Y)$
- ▶ Conditional independence: $P(X, Y|Z) = P(X|Z)P(Y|Z)$

Naïve Bayes: the objective

The objective is to maximize the joint probability of a set of labeled training documents $p(\mathbf{x}^{1:N}, y^{1:N})$, where N is the number of documents. This is known as the **maximum likelihood estimation**.

The goal of the training process is to find the weights θ that maximizes this likelihood.

$$\hat{\theta} = \operatorname{argmax}_{\theta} p(\mathbf{x}^{1:N}, y^{1:N}; \theta)$$

$$= \operatorname{argmax}_{\theta} \prod_{i=1}^N p(\mathbf{x}^{(i)}, y^{(i)}; \theta)$$

$$= \operatorname{argmax}_{\theta} \sum_{i=1}^N \log p(\mathbf{x}^{(i)}, y^{(i)}; \theta)$$

The notation $p(\mathbf{x}^{1:N}, y^{1:N}; \theta)$ indicate that θ is a parameter of the probability function. Symbols in bold indicate a vector of variables rather than a single variable.

“Independent and Identically Distributed”

<https://powcoder.com>

Assignment Project Exam Help

- ▶ “A collection of random variables is **independent and identically distributed** if each random variable has the same probability distribution as the others and all are mutually independent.” - from Wikipedia
<https://powcoder.com>
- ▶ Often shortened as *i.i.d*
- ▶ The basis on which we can break down the joint probability of all samples into the product of the probability of each sample

The generative story of Naïve Bayes

<https://powcoder.com>

The probability $p(\mathbf{x}^{1:N}, y^{1:N}; \theta)$ is defined through a **generative model**, an idealized random process that has generated the observed data. The algorithm that describes the generative model underlying the Naïve Bayes classifier with the parameters $\Theta = \{\mu, \phi\}$:

<https://powcoder.com>

Algorithm 1 Generative process for the Naïve Bayes classification model

for Instance $i \in \{1, 2, \dots, N\}$ **do**
 Draw the label $y^{(i)} \propto \text{Categorical}(\mu)$;
 Draw the word counts $\mathbf{x}^{(i)} | y^{(i)} \propto \text{Multinomial}(\phi^{(i)})$.
end for

Multinomial distribution

<https://powcoder.com>

Assignment Project Exam Help

$P_{Y|X}$ is a **multinomial** which a probabilistic distribution over vectors of non-negative counts. The probability mass function for this distribution is:

<https://powcoder.com>

Add WeChat powcoder

$$P_{\text{multi}}(\mathbf{x}, \phi) = B(\mathbf{x}) \prod_{j=1}^V \phi_j^{x_j}$$
$$B(\mathbf{x}) = \frac{\left(\sum_{j=1}^V x_j\right)!}{\prod_{j=1}^V (x_j!)}$$

Crucially, $B(\mathbf{x})$ is a multinomial coefficient that does not depend on ϕ , and can usually be ignored.

Parameter estimation

The generative story above allows us to decompose

$$\mathcal{L}(\theta) = \sum_{i=1}^N \log p(\mathbf{x}^{(i)}, y^{(i)}; \theta)$$

into

$$\begin{aligned} \mathcal{L}(\phi, \mu) &= \sum_{i=1}^N \log P_{mult}(\mathbf{x}^{(i)}; \phi_{y^{(i)}}) + \log P_{cat}(y^{(i)}; \mu) \\ &= \sum_{i=1}^N \log B(\mathbf{x}^{(i)}) + \sum_{j=1}^V x_j^{(i)} \log \phi_{y^{(i)}, j} + \sum_{i=1}^N \log \mu_{y^{(i)}} \end{aligned}$$

Maximum-likelihood estimation chooses ϕ and μ that maximize the log-likelihood of \mathcal{L} . Because we want these parameters to be probabilities, the solution must obey the following constraints:

$$\sum_{j=1}^V \phi_{y,j} = 1 \quad \forall y$$

Parameter Estimation

<https://powcoder.com>

After incorporating the the constraints by adding a set of Lagrange multipliers, we get this new objective (we focus on the $\phi_{y,j}$ for now):

Assignment Project Exam Help

Add WeChat powcoder

$$\ell(\Phi_y) = \sum_{i: y^{(i)}=y} \sum_{j=1}^V x_j^{(i)} \log \phi_{y,j} - \lambda \left(\sum_{j=1}^V \phi_{y,j} - 1 \right)$$

<https://powcoder.com>

Differentiating with respect to the parameters $\phi_{y,j}$ yields,

Add WeChat powcoder

$$\frac{\partial \ell(\Phi_y)}{\partial \phi_{y,j}} = \sum_{i: y^{(i)}=y} x_j^{(i)} / \phi_{y,j} - \lambda$$

Parameter Estimation

<https://powcoder.com>

There is a closed form solution to this which can be obtained by setting each element in the vector of derivatives equal to zero:

$$\lambda \phi_{y,j} = \sum_{i: y^{(i)}=y} x_j^{(i)}$$
$$\phi_{y,j} \propto \sum_{i: y^{(i)}=y} x_j^{(i)} = \sum_i \delta(y^{(i)} = y) x_j^{(i)} = \text{count}(y, j)$$

where $\delta(y^{(i)} = y)$ is an indicator function which returns one if $y^{(i)} = y$. The symbol \propto indicates that $\phi_{y,j}$ is proportional to the right-hand side of the equation

Parameter Estimation

<https://powcoder.com>

- Recall the constraint that ϕ_y is a vector of probabilities:
 $\sum_{j=1}^V \phi_{y,j} = 1$. We have an exact solution:

$$\phi_{y,j} = \frac{\text{count}(y,j)}{\sum_{j'=1}^V \text{count}(y,j')}$$

<https://powcoder.com>

- Similarly we can arrive at:

$$\mu_y = \frac{\text{count}(y)}{\sum_{y' \in K} \text{count}(y')}$$

As is often the case, the result of the mathematical derivation (that needs to be turned into code) is usually often much simpler:

Smoothing

<https://powcoder.com>

Laplace smoothing: add α

$$\phi_{y,j} = \frac{\alpha + \text{count}(y,j)}{\alpha + \sum_{j=1}^J \text{count}(y,j)}$$

The **bias - variance** trade-off:

- ▶ Unbiased classifiers may **overfit** the training data, yielding poor performance on unseen data.
- ▶ But if the smoothing is too large, the resulting classifier can **underfit** instead. In the limit of $\alpha \rightarrow \infty$, there is zero variance: you get the same classifier, regardless of the data.

How to determine the best α ?

Grid Search

<https://powcoder.com>

Setting hyperparameters with grid search. parameters and hyperparameters

- ▶ Try a set of values and find the one that maximizes the accuracy, but on which data set?
- ▶ The goal is to maximize the system on *unseen* data, so we shouldn't be doing it on training data. Instead, we should do it on a *development or tuning* set.
- ▶ We should also set aside another set called *test* set that you measure system performance on.
- ▶ If the data set is too small, use *cross-validation*

Prediction with Naïve Bayes

<https://powcoder.com>

$$\hat{y} = \operatorname{argmax} \log p(\mathbf{x}, y; \boldsymbol{\mu}, \boldsymbol{\phi})$$

$$= \operatorname{argmax} \log p(\mathbf{x}|y; \boldsymbol{\phi}) + \log p(y; \boldsymbol{\mu})$$

Plug in the distribution from the generative story, we get:

$$\log p(\mathbf{x}|y; \boldsymbol{\phi}) + \log p(y; \boldsymbol{\mu}) = \log \left[B(\mathbf{x}) \prod_{j=1}^V \phi_{y,j}^{x_j} \right] + \log \mu_y$$

$$\begin{aligned} &= \log B(\mathbf{x}) + \sum_{j=1}^V x_j \log(\phi_{y,j}) + \log \mu_y \\ &= \log B(\mathbf{x}) + \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}, y), \end{aligned}$$

where

$$\boldsymbol{\theta} = [\boldsymbol{\theta}^{(1)}; \boldsymbol{\theta}^{(2)}; \dots; \boldsymbol{\theta}^{(K)}]$$

$$\boldsymbol{\theta}^{(y)} = [\log \phi_{y,1}; \log \phi_{y,2}; \dots; \log \phi_{y,V}; \log \mu_y]$$

Relation between mathematical models and computer science tools

<https://powcoder.com>

- ▶ Mathematics provides the justification of why a model works the way it does, and computer science focuses on realizing it with efficient algorithms and appropriate data structures.
 - ▶ Computational algorithms often resort to caching to avoid repeated computation, thus making the computational implementation more efficient, e.g., Viterbi, Forward-Backward, CKY
- ▶ It's useful to think about where the mathematical justification ends and computational realization starts.
- ▶ The relation between a mathematical expression and its implementation in a programming language can be thought of as a translation process: it's not always word for word.

Advantages of Naïve Bayes

<https://powcoder.com>

Assignment Project Exam Help

With a joint likelihood objective, the estimated parameters of a Naïve Bayes model can be used for both classification ($p(y|x)$) and generation ($p(x|y)$).

<https://powcoder.com>

$$P(x, y) = P(x|y) \times P(y) = P(y|x) \times P(x)$$

Add WeChat powcoder

In practice, we rarely use Naïve Bayes for generation. It's mostly used as a classification model.

Problems of Naïve Bayes

<https://powcoder.com>

Let's say we want to include some subword units (e.g., morphemes).

Assignment Project Exam Help

$$\begin{aligned} P(\text{word} = \text{unfit}, \text{prefix} = \text{un-} | y) \\ &= P(\text{prefix} = \text{un-} | \text{word} = \text{unfit}, y) \times P(\text{word} = \text{unfit} | y) \\ &= 1 \times P(\text{word} = \text{unfit} | y) \end{aligned}$$

Add WeChat powcoder

<https://powcoder.com>

If we assume conditional independence,

$$\begin{aligned} P(\text{word} = \text{unfit}, \text{prefix} = \text{un-} | y) \\ \approx P(\text{word} = \text{unfit} | y) \times P(\text{prefix} = \text{un-} | y) \end{aligned}$$

Add WeChat powcoder

Since $P(\text{word} = \text{unfit} | y) \geq P(\text{word} = \text{unfit} | y) \times P(\text{prefix} = \text{un-} | y)$, conditional independence under-estimates the true probabilities of conjunction of positively correlated features.