

## Linear models: Recap

<https://powcoder.com>

Linear models:

- Perceptron

Assignment Project Exam Help

Add WeChat powcoder

- Naïve Bayes:

<https://powcoder.com>

$$\log P(y|\mathbf{x}; \theta) = \log P(\mathbf{x}|y; \phi) + \log P(y; \mathbf{u}) = \log B(\mathbf{x}) + \theta \cdot \mathbf{f}(\mathbf{x}, y)$$

Add WeChat powcoder

- Logistic Regression

$$\log P(y|\mathbf{x}; \theta) = \theta \cdot \mathbf{f}(\mathbf{x}, y) - \log \sum_{y' \in \mathcal{Y}} \exp \theta \cdot \mathbf{f}(\mathbf{x}, y')$$

## Features and weights in linear models: Recap

- Feature representation:  $f(\mathbf{x}, y)$

$$f(\mathbf{x}, y=1) = [x; \underbrace{0, 0, \dots, 0}_{(K-1) \times V}]$$

$$f(\mathbf{x}, y=2) = [\underbrace{0, 0, \dots, 0}_V; x; \underbrace{0, 0, \dots, 0}_{(K-2) \times V}]$$

$$f(\mathbf{x}, y=K) = [\underbrace{0, 0, \dots, 0}_{(K-1) \times V}; x]$$

- Weights:  $\theta$

$$\theta = [\underbrace{\theta_1; \theta_2; \dots; \theta_V}_{y=1}; \underbrace{\theta_1; \theta_2; \dots; \theta_V}_{y=2}; \dots; \underbrace{\theta_1; \theta_2; \dots; \theta_V}_{y=K}]$$

## Rearranging the features and weights

<https://powcoder.com>

- Represent the features  $\mathbf{x}$  as a *column* vector of length  $V$ , and represent the weights as a  $\Theta$  as  $K \times V$  matrix

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_V \end{bmatrix} \quad \Theta = \begin{matrix} & \begin{matrix} x_1 & x_2 & \cdots & x_V \end{matrix} \\ \begin{matrix} y = 1 \\ y = 2 \\ \vdots \\ y = K \end{matrix} & \begin{bmatrix} \theta_{1,1} & \theta_{1,2} & \cdots & \theta_{1,V} \\ \theta_{2,1} & \theta_{2,2} & \cdots & \theta_{2,V} \\ \vdots & \vdots & \ddots & \vdots \\ \theta_{K,1} & \theta_{K,2} & \cdots & \theta_{K,V} \end{bmatrix} \end{matrix}$$

- What is  $\Theta \mathbf{x}$ ?

Scores for each class

<https://powcoder.com>

## Assignment Project Exam Help

- Verify that  $\psi_1, \psi_2, \dots, \psi_K$  correspond to the scores for each class

<https://powcoder.com>

$$\Psi = \Theta \mathbf{x} = \begin{bmatrix} \theta_1 \cdot \mathbf{x} = \psi_1 \\ \theta_2 \cdot \mathbf{x} = \psi_2 \\ \vdots \\ \theta_K \cdot \mathbf{x} = \psi_K \end{bmatrix}$$

Add WeChat powcoder

## Implementation in Pytorch

<https://powcoder.com>

Assignment Project Exam Help

```
In [48]: weights = torch.randn(3,9)
print(weights)
input = torch.randn(9)
print(input)
output = torch.matmul(weights, input)
print(output)
softmax = nn.Softmax(dim=0)
probs = softmax(output)
print(probs)

tensor([[ -0.3518, -0.3291,  1.6937, -0.9947, -0.1390, -0.7788, -0.0252, -0.1557,
          -1.2138],
        [-1.2000,  1.3527,  1.9529, -1.3182,  0.1101, -0.7105, -0.4409,  0.9753,
          -0.0821],
        [-0.5561, -0.1114,  0.4833,  0.9997,  0.5840, -1.4133,  1.1353,  0.3069,
          -0.0584]])
tensor([ 0.0148,  0.0565, -0.6462, -0.0155, -0.5532, -0.8514, -0.1339,  0.5056,
         0.6025])
tensor([-1.1695, -0.1363,  0.5428])
tensor([0.1069, 0.3005, 0.5926])
```

<https://powcoder.com>

Add WeChat powcoder

Digression: Matrix multiplication

<https://powcoder.com>

## Assignment Project Exam Help

- ▶ Matrix with  $m$  rows and  $n$  columns:

[Add WeChat powcoder](https://powcoder.com)  
 $A \in \mathbb{R}^{m \times n}, B \in \mathbb{R}^{n \times p}, C = AB \in \mathbb{R}^{m \times p}$

where  $C_{ij} = \sum_{k=1}^n A_{ik}B_{kj}$  <https://powcoder.com>

- ▶ Example:

[Add WeChat powcoder](https://powcoder.com)

$$\begin{bmatrix} 2 & 3 \\ 1 & 2 \end{bmatrix} \times \begin{bmatrix} 1 & 0 & 2 \\ 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 2 & 3 & 7 \\ 1 & 2 & 4 \end{bmatrix}$$

## Digression: 3-D matrix multiplication

```
► In [27]: import torch
input = torch.randint(5, (2, 3, 4))
print(input)
mat2 = torch.randint(5, (2, 4, 3))
print(mat2)
out = torch.bmm(input, mat2)
print(out)
```

```
tensor([[[0, 1, 3, 2],
         [3, 1, 4, 1],
         [0, 1, 2, 6]]],
```

```
       [[3, 2, 3, 4],
        [4, 3, 2, 4],
        [3, 1, 3, 3]])
tensor([[[0, 2, 0],
         [4, 4, 2],
         [2, 2, 4],
         [1, 1, 4]]],
```

```
       [[3, 0, 4],
        [3, 0, 1],
        [0, 0, 4],
        [2, 4, 2]])
tensor([[[12, 12, 22],
         [13, 19, 22],
         [ 8,  8, 10]],
```

```
       [[23, 16, 34],
        [29, 16, 35],
        [15, 12, 26]])])
```

Tensor shape: (batch-size, sentence-length, embedding size)

## SoftMax

- ▶ SoftMax, also known as normalized exponential function.

Assignment Project Exam Help

$$\text{SoftMax}_i(\psi) = \frac{\exp \psi_i}{\sum_j^K \exp \psi_j}$$

Assignment Project Exam Help  
for  $i = 1, 2, \dots, K$

- ▶ Applying SoftMax turns the scores into a probabilistic distribution:

Add WeChat powcoder

$$\text{SoftMax}(\Psi) = \begin{bmatrix} P(y = 1) \\ P(y = 2) \\ \dots \\ P(y = K) \end{bmatrix}$$

- ▶ Verify this is exactly logistic regression



Logistic regression as a neural network


<https://powcoder.com>

Assignment Project Exam Help

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder




$$\mathbf{y} = \text{SoftMax}(\mathbf{\Theta} \mathbf{x})$$

$$V = 5 \quad K = 3$$

## Going deep

- There is no reason why we can't add layers in the middle




$$\mathbf{z} = \sigma(\Theta_1 \mathbf{x})$$

$$\mathbf{y} = \text{SoftMax}(\Theta_2 \mathbf{z})$$

## Going even deeper

- There is no reason why we can't add layers in the middle



$$\mathbf{z}_1 = \sigma(\Theta_1 \mathbf{x})$$

$$\mathbf{z}_2 = \sigma(\Theta_2 \mathbf{z}_1)$$

$$\mathbf{y} = \text{SoftMax}(\Theta_3 \mathbf{z}_2)$$

- But why?

## Non-linear classification

Linear models like Logistic regression can map data into a high-dimensional vector space, and they are expressive enough and work well for many NLP problems, why do we need more complex non-linear models?

- ▶ There have been rapid advances in deep learning, a family of nonlinear methods that learn complex functions of the input through multiple layers of computation.
- ▶ Deep learning facilitates the incorporation of **word embeddings**, which are dense vector representations of words, that can be learned from massive amounts of unlabeled data.
  - ▶ It has evolved from early static embeddings (e.g., Word2vec, Glove) to recent dynamic embeddings (ELMO, BERT, XLNet)
- ▶ Rapid advances in specialized hardware called graphic processing units (GPUs). Many deep learning models can be implemented efficiently on GPUs.

## Feedforward Neural networks: an intuitive justification

<https://powcoder.com>

- ▶ In image classification, instead of using the input (pixels) to predict the image type directly, you can imagine a scenario that you can predict the shapes of parts of an image, mouth, hand, ear.
- ▶ In text processing, we can imagine a similar scenario. Let's say we want to classify movie reviews (or movies themselves) into a label set of {Good, Bad, OK}. Instead predicting these labels directly, we first predict a set of composite features such as the story, acting, soundtrack, cinematography, etc. from raw input (words in the text).

## Face Recognition

<https://powcoder.com>

Assignment Project Exam Help

Assignment Project Exam Help



Layer 3

<https://powcoder.com>

Add WeChat powcoder



Layer 2



Layer 1

## Feedforward neural networks

<https://powcoder.com>

Formally, this is what we do:

- ▶ **Use the text  $x$  to predict the features  $z$ .** Specifically, train a logistic regression classifier to compute  $P(z_k|x)$ , for each  $k \in \{1, 2, \dots, K_z\}$
- ▶ **Use the features  $z$  to predict the label  $y$ .** Train a logistic regression classifier to compute  $P(y|z)$ .  $z$  is unknown or hidden, so we will use the  $P(z|x)$  as the features.

Caveat: it's easy to demonstrate what this is what the model does for image processing, but it's hard to show this is what's actually going on in language processing. Interpretability is a major issue in neural models for language processing.

The hidden layer: computing the composite features

<https://powcoder.com>

- If we assume each  $z_k$  is binary, that is,  $z_k \in \{0, 1\}$ , then  $P(z_k|\mathbf{x})$  can be modeled with binary logistic regression

$$P(z_k = 1|\mathbf{x}; \Theta^{(x \rightarrow z)}) = \sigma(\theta_k^{x \rightarrow z} \cdot \mathbf{x}) = (1 + \exp(-\theta_k^{x \rightarrow z} \cdot \mathbf{x}))^{-1}$$

- The weight matrix  $\Theta^{(x \rightarrow z)} \in \mathbb{R}^{k_z \times V}$  is constructed by stacking (not concatenating, as in linear models) the weight vectors for each  $z_k$ ,  
 $\Theta^{(x \rightarrow z)} = [\theta_1^{x \rightarrow z}, \theta_2^{x \rightarrow z}, \dots, \theta_{K_z}^{x \rightarrow z}]^T$
- We assume an offset/bias term is included in  $\mathbf{x}$  and its parameter is included in each  $\theta_k^{x \rightarrow z}$

Notations:  $\Theta^{(x \rightarrow z)} \in \mathbb{R}^{k_z \times V}$  is a real number matrix with a dimension of  $k_z$  rows and  $V$  columns



## The output layer

<https://powcoder.com>

- ▶ The output layer is computed by the multiclass logistic regression probability

$$P(y = j | \mathbf{z}; \Theta^{(z \rightarrow y)}, \mathbf{b}) = \frac{\exp(\theta_j^{(z \rightarrow y)} \cdot \mathbf{z} + b_j)}{\sum_{j' \in \mathcal{Y}} \exp(\theta_{j'}^{(z \rightarrow y)} \cdot \mathbf{z} + b_{j'})}$$

<https://powcoder.com>

- ▶ The weight matrix  $\Theta^{(z \rightarrow y)} \in \mathbb{R}^{k_y \times k_z}$  again is constructed by stacking weight vectors for each  $y$ :

$$\Theta^{(z \rightarrow y)} = \begin{bmatrix} \theta_1^{z \rightarrow y} & \theta_2^{z \rightarrow y} & \dots & \theta_{K_y}^{z \rightarrow y} \end{bmatrix}^\top$$

- ▶ The vector of probabilities over each possible value of  $y$  is denoted:

$$P(\mathbf{y} | \mathbf{z}; \Theta^{(z \rightarrow y)}, \mathbf{b}) = \text{SoftMax}(\Theta^{(z \rightarrow y)} \mathbf{z} + \mathbf{b})$$

## Activation functions

- Sigmoid: The range of sigmoid function is  $(0, 1)$ .

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

- Tanh: The range of the tanh activation function is  $(-1, 1)$

$$\tanh(x) = \frac{e^{2x} - 1}{e^{2x} + 1}$$

- ReLU: The rectified linear unit (ReLU) is zero for negative inputs, and linear for positive inputs

$$ReLU(x) = \max(x, 0) = \begin{cases} 0 & x < 0 \\ x & \text{otherwise} \end{cases}$$

Sigmoid and tanh are sometimes described as **squashing functions**.

## Activation functions in Pytorch

<https://powcoder.com>

```
from torch import nn
import torch

input = torch.randn(4)
sigmoid = nn.Sigmoid()
output = sigmoid(input)

tanh = nn.Tanh()
output = tanh(input)

relu = nn.ReLU()
output = relu(input)
```

## Output and loss functions

<https://powcoder.com>  
In a multi-class classification setting, a softmax output produces a probabilistic distribution over possible labels. It works well together with negative conditional likelihood (just like logistic regression)

<https://powcoder.com>  
[Add WeChat powcoder](#)

$$-\mathcal{L} = - \sum_{i=1}^N \log P(y^{(i)} | \mathbf{x}^{(i)}; \Theta)$$

<https://powcoder.com>  
or **cross entropy loss**:

[Add WeChat powcoder](#)

$$\tilde{y}_j \triangleq \Pr(y = j | \mathbf{x}^{(i)}; \Theta)$$
$$-\mathcal{L} = - \sum_{i=1}^N \mathbf{e}_{y^{(i)}} \cdot \log \tilde{\mathbf{y}}$$

where  $\mathbf{e}_{y^{(i)}}$  is a **one-hot vector** of zeros with a value of one at the position  $y^{(i)}$

## Output and loss function

<https://powcoder.com>

### Assignment Project Exam Help

- ▶ There are alternatives to SoftMax and cross-entropy loss, just as there are alternatives in linear models.
- ▶ Pairing an affine transformation (remember perceptron) with a margin loss:

<https://powcoder.com>

$$\Psi(y; \mathbf{x}^{(i)}, \Theta) = \theta_y^{(z \rightarrow y)} \mathbf{z} + b_y$$

Add WeChat powcoder

$$\ell_{\text{MARGIN}}(\Theta; \mathbf{x}^{(i)}, y^{(i)}) = \max_{y \neq y^{(i)}} \left( 1 + \Psi(y; \mathbf{x}^{(i)}, \Theta) - \Psi(y^{(i)}; \mathbf{x}^{(i)}, \Theta) \right)_+$$

## Inputs and Lookup layers

<https://powcoder.com>

- ▶ Assuming a bag-of-words model, when the input  $\mathbf{x}$  is the count of each word  $x_j$  (This can be generalized to feature count).
- ▶ To compute the hidden unit  $z_k$ :

$$z_k = \sum_{j=1}^V \theta_{j,k}^{x \rightarrow z} x_j$$

- ▶ This text representation is particularly suited for feedforward networks.
- ▶ The connections from word  $j$  to each of the hidden units  $z_k$  form a vector  $\theta_j^{(x \rightarrow z)}$  is sometimes described as the embedding of word  $j$ . Word embeddings can be learned from unlabeled data, using techniques such as Word2Vec and GLOVE.

## Alternative text representations

<https://powcoder.com>

- ▶ Alternatively, a text can be represented as a sequence of word tokens  $w_1, w_2, w_3, \dots, w_M$ . This view is useful for models such as **Convolutional Neural Networks**, or **ConvNets**, which processes text as a sequence.
- ▶ Each word token  $w_m$  is represented as a one-hot vector  $\mathbf{e}_{w_m}$ , with dimension  $V$ . The complete document can be represented by the horizontal concatenation of these one-hot vectors:  $\mathbf{W} = [\mathbf{e}_{w_1}, \mathbf{e}_{w_2}, \dots, \mathbf{e}_{w_M}] \in R^{V \times M}$
- ▶ To show that this is equivalent to the bag-of-words model, we can recover the word count from the matrix-vector product  $\mathbf{W}[1, 1, \dots, 1]^T \in R^V$ .
- ▶ The matrix product  $\Theta^{x \rightarrow z} \mathbf{W} \in R^{k_z \times M}$  contains horizontally concatenated embeddings of each word in the document.